

# Analysis of RNA-Seq Data with R/Bioconductor

...

Thomas Girke

December 8, 2012

Overview

RNA-Seq Analysis

Aligning Short Reads

Viewing Results in IGV Genome Browser

# Outline

## Overview

## RNA-Seq Analysis

### Aligning Short Reads

## Viewing Results in IGV Genome Browser

# Packages for RNA-Seq Analysis in R

- GenomicRanges [Link](#): high-level infrastructure for range data
- Rsamtools [Link](#): BAM support
- rtracklayer [Link](#): Annotation imports, interface to online genome browsers
- DESeq [Link](#): RNA-Seq DEG analysis
- edgeR [Link](#): RNA-Seq DEG analysis
- DEXSeq [Link](#): RNA-Seq Exon analysis

# RNA-Seq versus DGE

## RNA-seq




Sequencing



1. Alternative splicing
2. Limited expression profiling
3. SNP detection
4. Many other applications

## DGE

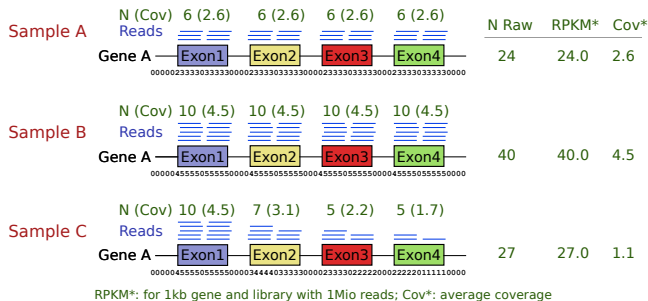


Sequencing



1. Expression profiling
- more appropriate for many biosamples

# Identification of Differentially Expressed Genes



Normalization often by library size.

# RNA-Seq Analysis Workflow

- Read mapping
- Counting reads overlapping with genes
- Analysis of differentially expressed genes (DEGs)
- Clustering of co-expressed genes
- Gene set/GO term enrichment analysis

# Outline

Overview

RNA-Seq Analysis

Aligning Short Reads

Viewing Results in IGV Genome Browser



# Data Sets and Experimental Variables

- To make the following sample code work, please download and unpack the sample data [Link](#) in the directory of your current R session.
- It contains four simplified alignment files from RNA-Seq experiment SRA023501 [Link](#) and a shortened GFF to allow fast analysis on a laptop.
- The alignments were created by aligning the reads with Bowtie against the Arabidopsis reference genome.
- **Note:** usually, the aligned reads would be stored in BAM format and then imported into R with the `readBamGappedAlignments` function (see below)!

This information could be imported from an external targets file

```
> targets <- read.delim("./data/targets.txt")  
> targets
```

	Samples	Factor	Fastq
1	AP3_f14	AP3	SRR064154.fastq
2	AP3_f14	AP3	SRR064155.fastq
3	T1_f14	TRL	SRR064166.fastq
4	T1_f14	TRL	SRR064167.fastq

# Outline

Overview

RNA-Seq Analysis

Aligning Short Reads

Viewing Results in IGV Genome Browser

# Align Reads and Output Indexed Bam Files

Note: this step requires the command-line tool bowtie2 [Link](#). If it is not available on a system then one can skip this mapping step and download the pre-generated Bam files from here: [Link](#)

```
> library(Rsamtools)
> dir.create("results") # Note: all output data will be written to directory 'results'
> system("bowtie2-build ./data/tair10chr.fasta ./data/tair10chr.fasta") # Build indexed reference genome
> targets <- read.delim("./data/targets.txt") # Import experiment design information
> targets
> input <- paste("./data/", targets[,3], sep="")
> output <- paste("./results/", targets[,3], ".sam", sep="")
> reference <- "./results/tair10chr.fasta"
> for(i in seq(along=targets[,3])) {
+   command <- paste("bowtie2 -x ./data/tair10chr.fasta -U", input[i], "-S", output[i])
+   system(command)
+   asBam(file=output[i], destination=gsub(".sam", "", output[i]), overwrite=TRUE, indexDestination=TRUE)
+   unlink(output[i])
+ }
```

# Import Annotation Data from GFF

## Annotation data from GFF

```
> library(rtracklayer); library(GenomicRanges); library(Rsamtools)
> gff <- import.gff("./data/TAIR10_GFF3_trunc.gff", asRangedData=FALSE)
> seqlengths(gff) <- end(ranges(gff[which(elementMetadata(gff)[,"type"]=="chromosome"),]))
> subgene_index <- which(elementMetadata(gff)[,"type"] == "gene")
> gffsub <- gff[subgene_index,] # Returns only gene ranges
> strand(gffsub) <- "*" # For strand insensitive analysis
> gffsub[1:4,1:2]
```

GRanges with 4 ranges and 2 metadata columns:

	seqnames	ranges	strand	source	type
	<Rle>	<IRanges>	<Rle>	<factor>	<factor>
[1]	Chr1	[ 3631, 5899]	*	TAIR10	gene
[2]	Chr1	[ 5928, 8737]	*	TAIR10	gene
[3]	Chr1	[11649, 13714]	*	TAIR10	gene
[4]	Chr1	[23146, 31227]	*	TAIR10	gene

---

seqlengths:

Chr1	Chr2	Chr3	Chr4	Chr5	ChrC	ChrM
100000	100000	100000	100000	100000	100000	100000

```
> ids <- as.character(elementMetadata(gffsub)[, "group"])
> gffsub <- split(gffsub, ids) # Coerce to GRangesList
```

# Read Counting per Annotation Range - Old

## Number of reads overlapping gene ranges

```
> samples <- as.character(targets$Fastq)
> samplespath <- paste("./results/", samples, ".bam", sep="")
> names(samplespath) <- samples
> countDF <- data.frame(row.names=ids)
> for(i in samplespath) {
+   aligns <- readBamGappedAlignments(i) # Substitute next two lines with this one.
+   counts <- countOverlaps(gffsub, aligns)
+   countDF <- cbind(countDF, counts)
+ }
> colnames(countDF) <- samples
> rownames(countDF) <- gsub(".*=", "", rownames(countDF))
> countDF[1:4,]
```

	SRR064154.fastq	SRR064155.fastq	SRR064166.fastq	SRR064167.fastq
AT1G01010	52	26	60	75
AT1G01020	149	79	84	66
AT1G01030	5	1	13	14
AT1G01040	492	355	303	365

```
> write.table(countDF, "./results/countDF", quote=FALSE, sep="\t", col.names = NA)
> countDF <- read.table("./results/countDF")
```

# Read Counting per Annotation Range - New

The `summarizeOverlaps` function from the `GenomicRanges` is easier to use and provides more options. See [here](#)

[Link](#) for details.

```
> bfl <- BamFileList(samplespath, index=character())
> countDF2 <- summarizeOverlaps(gffsub, bfl, mode="Union", ignore.strand=TRUE)
> countDF2 <- assays(countDF2)$counts
> rownames(countDF2) <- gsub(".*=", "", rownames(countDF2))
> countDF2[1:4,]
```

	SRR064154.fastq	SRR064155.fastq	SRR064166.fastq	SRR064167.fastq
AT1G01010	52	26	60	75
AT1G01020	149	79	84	66
AT1G01030	5	1	13	14
AT1G01040	480	346	282	336

# Simple RPKM Normalization

RPKM: reads per kilobase of exon model per million mapped reads


```
> returnRPKM <- function(counts, gffsub) {  
+   geneLengthsInKB <- sum(width(gffsub))/1000 # Number of bases per exonRanges element in kbp  
+   millionsMapped <- sum(counts)/1e+06 # Factor for converting to million of mapped reads.  
+   rpm <- counts/millionsMapped # RPK: reads per kilobase of exon model.  
+   rpkm <- rpm/geneLengthsInKB # RPKM: reads per kilobase of exon model per million mapped reads.  
+   return(rpkm)  
+ }  
> countDFrpkm <- apply(countDF, 2, function(x) returnRPKM(counts=x, gffsub=gffsub))  
> countDFrpkm[1:4,]
```

	SRR064154.fastq	SRR064155.fastq	SRR064166.fastq	SRR064167.fastq
AT1G01010	38.961967	18.3057160	287.8127	271.20512
AT1G01020	90.147145	44.9126457	325.3615	192.71199
AT1G01030	4.114449	0.7732458	68.4867	55.59924
AT1G01040	103.494741	70.1709530	408.0534	370.54857

# QC Check

QC check by computing a sample correlating matrix and plotting it as a tree

```
> d <- cor(countDF, method="spearman")  
> plot(hclust(dist(1-d))) # Sample tree
```





# Identify DEGs with Simple Fold Change Method

## Compute mean values for replicates

```
> source("http://faculty.ucr.edu/~tgirke/Documents/R_BioCond/My_R_Scripts/colAg.R")
> countDfrpkm_mean <- colAg(myMA=countDfrpkm, group=c(1,1,2,2), myfct=mean)
> countDfrpkm_mean[1:4,]
```

	SRR064154.fastq_SRR064155.fastq	SRR064166.fastq_SRR064167.fastq
AT1G01010	28.633841	279.50891
AT1G01020	67.529895	259.03674
AT1G01030	2.443848	62.04297
AT1G01040	86.832847	389.30097

## Log2 fold changes

```
> countDfrpkm_mean <- cbind(countDfrpkm_mean, log2ratio=log2(countDfrpkm_mean[,2]/countDfrpkm_mean[,1]))
> countDfrpkm_mean <- countDfrpkm_mean[is.finite(countDfrpkm_mean[,3]), ]
> degs2fold <- countDfrpkm_mean[countDfrpkm_mean[,3] >= 1 | countDfrpkm_mean[,3] <= -1,]
> degs2fold[1:4,]
```

	SRR064154.fastq_SRR064155.fastq	SRR064166.fastq_SRR064167.fastq
AT1G01010	28.633841	279.50891
AT1G01020	67.529895	259.03674
AT1G01030	2.443848	62.04297
AT1G01040	86.832847	389.30097

	log2ratio
AT1G01010	3.287101
AT1G01020	1.939559
AT1G01030	4.666042
AT1G01040	2.164573

```
> write.table(degs2fold, "./results/degs2fold", quote=FALSE, sep="\t", col.names = NA)
> degs2fold <- read.table("./results/degs2fold")
```

# Identify DEGs with DESeq Library

Raw count data are expected here!

```
> library(DESeq)
> countDF <- read.table("./results/countDF")
> conds <- targets$Factor
> cds <- newCountDataSet(countDF, conds) # Creates object of class CountDataSet derived from eSet class
> counts(cds)[1:4, ] # CountDataSet has similar accessor methods as eSet class.
```

	SRR064154.fastq	SRR064155.fastq	SRR064166.fastq	SRR064167.fastq
AT1G01010	52	26	60	75
AT1G01020	149	79	84	66
AT1G01030	5	1	13	14
AT1G01040	492	355	303	365

```
> cds <- estimateSizeFactors(cds) # Estimates library size factors from count data. Alternatively, one can
> cds <- estimateDispersions(cds) # Estimates the variance within replicates
> res <- nbinomTest(cds, "AP3", "TRL") # Calls DEGs with nbinomTest
> res <- na.omit(res)
> res2fold <- res[res$log2FoldChange >= 1 | res$log2FoldChange <= -1,]
> res2foldpadj <- res2fold$res2fold$padj <= 0.05, ]
> res2foldpadj[1:4,1:8]
```

	id	baseMean	baseMeanA	baseMeanB	foldChange	log2FoldChange
6	AT1G01050	602.73819	270.09411	935.3823	3.463172	1.792094
7	AT1G01060	302.65623	168.64177	436.6707	2.589339	1.372584
8	AT1G01070	30.06197	5.62154	54.5024	9.695281	3.277283
12	AT1G01100	6353.10140	4226.29888	8479.9039	2.006461	1.004653
	pval	padj				
6	5.840333e-13	6.891592e-12				
7	2.433388e-06	1.511262e-05				
8	4.413539e-05	2.003068e-04				
12	4.243839e-07	2.782072e-06				

# Identify DEGs with edgeR Library

Raw count data are expected here!

```
> library(edgeR)
> countDF <- read.table("./results/countDF")
> y <- DGEList(counts=countDF, group=conds) # Constructs DGEList object
> y <- estimateCommonDisp(y) # Estimates common dispersion
> y <- estimateTagwiseDisp(y) # Estimates tagwise dispersion
> et <- exactTest(y, pair=c("AP3", "TRL")) # Computes exact test for the negative binomial distribution.
> topTags(et, n=4)
```

Comparison of groups: TRL-AP3

	logFC	logCPM	PValue	FDR
AT4G00050	5.433513	10.35929	7.946221e-49	1.144256e-46
AT1G01050	4.203520	12.08703	1.013791e-45	7.299293e-44
AT3G01120	3.865095	14.07493	5.930030e-41	2.846415e-39
ATMG00030	-3.745470	10.89011	2.594715e-34	9.340974e-33

```
> edge <- as.data.frame(topTags(et, n=50000))
> edge2fold <- edge[edge$logFC >= 1 | edge$logFC <= -1,]
> edge2foldpadj <- edge2fold[edge2fold$FDR <= 0.01, ]
```


# Merge Results and Compute Overlaps Among Methods

```
> bothDF <- merge(res, countDFrpkm_mean, by.x=1, by.y=0, all=TRUE); bothDF <- na.omit(bothDF)
> cor(bothDF[, "log2FoldChange"], bothDF[, "log2ratio"], method="spearman")
```

```
[1] 0.9990604
```

```
> source("http://faculty.ucr.edu/~tgirke/Documents/R_BioCond/My_R_Scripts/overLapper.R")
> setlist <- list(edgeR=row.names(edge2foldpadj), DESeq=as.character(res2foldpadj[,1]), RPKM=row.names(degs2f))
> OList <- overLapper(setlist=setlist, sep="_", type="vennsets")
> counts <- sapply(OList$Venn_List, length)
> vennPlot(counts=counts)
```

Venn Diagram



Unique objects: All = 84; S1 = 62; S2 = 37; S3 = 72

# Enrichment of GO Terms in DEG Sets

## GO Term Enrichment Analysis

```
> library(GOstats); library(GO.db); library(ath1121501.db)
> geneUniverse <- rownames(countDF)
> geneSample <- res2foldpadj[,1]
> params <- new("GOHyperGParams", geneIds = geneSample, universeGeneIds = geneUniverse,
+   annotation="ath1121501", ontology = "MF", pvalueCutoff = 0.5,
+   conditional = FALSE, testDirection = "over")
> hgOver <- hyperGTest(params)
> summary(hgOver)[1:4,]
```

	GOMFID	Pvalue	OddsRatio	ExpCount	Count	Size	
1	GO:0008324	0.004436386	15.6	2.303797	6	7	
2	GO:0015075	0.004436386	15.6	2.303797	6	7	
3	GO:0015077	0.004436386	15.6	2.303797	6	7	
4	GO:0015078	0.004436386	15.6	2.303797	6	7	

		Term
1	cation transmembrane transporter activity	
2	ion transmembrane transporter activity	
3	monovalent inorganic cation transmembrane transporter activity	
4	hydrogen ion transmembrane transporter activity	

```
> htmlReport(hgOver, file = "data/MyhyperGresult.html")
```

# Outline

Overview

RNA-Seq Analysis

Aligning Short Reads

Viewing Results in IGV Genome Browser

# Inspect Results in IGV

## View results in IGV

- 1 Download and open IGV [Link](#)
- 2 Select in menu in top left corner A. thaliana (TAIR10)
- 3 Upload the following indexed/sorted Bam files with File -> Load from URL...


[http://faculty.ucr.edu/~tgirke/HTML\\_Presentations/Manuals/Workshop\\_Dec\\_6\\_10\\_2012/Rrnaseq/results/SRR064154.fastq](http://faculty.ucr.edu/~tgirke/HTML_Presentations/Manuals/Workshop_Dec_6_10_2012/Rrnaseq/results/SRR064154.fastq)

[http://faculty.ucr.edu/~tgirke/HTML\\_Presentations/Manuals/Workshop\\_Dec\\_6\\_10\\_2012/Rrnaseq/results/SRR064155.fastq](http://faculty.ucr.edu/~tgirke/HTML_Presentations/Manuals/Workshop_Dec_6_10_2012/Rrnaseq/results/SRR064155.fastq)

[http://faculty.ucr.edu/~tgirke/HTML\\_Presentations/Manuals/Workshop\\_Dec\\_6\\_10\\_2012/Rrnaseq/results/SRR064166.fastq](http://faculty.ucr.edu/~tgirke/HTML_Presentations/Manuals/Workshop_Dec_6_10_2012/Rrnaseq/results/SRR064166.fastq)

[http://faculty.ucr.edu/~tgirke/HTML\\_Presentations/Manuals/Workshop\\_Dec\\_6\\_10\\_2012/Rrnaseq/results/SRR064167.fastq](http://faculty.ucr.edu/~tgirke/HTML_Presentations/Manuals/Workshop_Dec_6_10_2012/Rrnaseq/results/SRR064167.fastq)

- 4 To view area of interest, enter its coordinates Chr1:49,457-51,457 in position menu on top.



# Analysis of Differential Exon Usage with DEXSeq

## Number of reads overlapping gene ranges

```
> source("data/Fct/gffexonDEXSeq.R")
> gffexonDEXSeq <- exons2DEXSeq(gff=gff)
> ids <- as.character(elementMetadata(gffexonDEXSeq)[, "ids"])
> countDFdex <- data.frame(row.names=ids)
> for(i in samplespath) {
+   aligns <- readBamGappedAlignments(i) # Substitute next two lines with this one.
+   counts <- countOverlaps(gffexonDEXSeq, aligns)
+   countDFdex <- cbind(countDFdex, counts)
+ }
> colnames(countDFdex) <- samples
> countDFdex[1:4,1:2]
```

SRR064154.fastq

Parent=AT1G01010:E001__Chr1_3631_3913+_Parent=AT1G01010.1	2
Parent=AT1G01010:E002__Chr1_3996_4276+_Parent=AT1G01010.1	2
Parent=AT1G01010:E003__Chr1_4486_4605+_Parent=AT1G01010.1	3
Parent=AT1G01010:E004__Chr1_4706_5095+_Parent=AT1G01010.1	6

SRR064155.fastq

Parent=AT1G01010:E001__Chr1_3631_3913+_Parent=AT1G01010.1	4
Parent=AT1G01010:E002__Chr1_3996_4276+_Parent=AT1G01010.1	1
Parent=AT1G01010:E003__Chr1_4486_4605+_Parent=AT1G01010.1	3
Parent=AT1G01010:E004__Chr1_4706_5095+_Parent=AT1G01010.1	1

```
> write.table(countDFdex, "./results/countDFdex", quote=FALSE, sep="\t", col.names = NA)
> countDFdex <- read.table("./results/countDFdex")
```



# Analysis of Differential Exon Usage with DEXSeq

## Identify genes with differential exon usage


```
> library(DEXSeq)
> samples <- as.character(targets$Factor); names(samples) <- targets$Fastq
> countDFdex[is.na(countDFdex)] <- 0
> ## Construct ExonCountSet from scratch
> exset <- newExonCountSet2(countDF=countDFdex) # fData(exset)[1:4,]
> ## Performs normalization
> exset <- estimateSizeFactors(exset)
> ## Evaluate variance of the data by estimating dispersion using Cox-Reid (CR) likelihood estimation
> exset <- estimateDispersions(exset)
> ## Fits dispersion-mean relation to the individual CR dispersion values
> exset <- fitDispersionFunction(exset)
> ## Performs Chi-squared test on each exon and Benjmini-Hochberg p-value adjustment for mutliple testing
> exset <- testForDEU(exset)
> ## Estimates fold changes of exons
> exset <- estimatelog2FoldChanges(exset)
> ## Obtain results in data frame
> deuDF <- DEUresultTable(exset)
> ## Count number of genes with differential exon usage
> table(tapply(deuDF$padjust < 0.01, geneIDs(exset), any))

FALSE TRUE
    20    1
```

## DEXSeq Plots

Sample plot showing fitted expression of exons

```
> plotDEXSeq(exset, "Parent=AT1G01100", displayTranscripts=TRUE, expression=TRUE, legend=TRUE)
> ## Generate many plots and write them to results directory
> mygeneIDs <- unique(as.character(na.omit(deuDF[deuDF$geneID %in% unique(deuDF$geneID),, "geneID"])))
> DEXSeqHTML(exset, geneIDs=mygeneIDs, path="results", file="DEU.html")
```



# Session Information

```
> sessionInfo()

R version 2.15.2 (2012-10-26)
Platform: x86_64-apple-darwin9.8.0/x86_64 (64-bit)
locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

attached base packages:
[1] stats      graphics  grDevices  utils      datasets  methods   base

other attached packages:
[1] DEXSeq_1.4.0          xtable_1.7-0          ath1121501.db_2.8.0
[4] org.At.tair.db_2.8.0 GO.db_2.8.0           GOstats_2.24.0
[7] RSQLite_0.11.2       DBI_0.2-5             graph_1.36.1
[10] Category_2.24.0      AnnotationDbi_1.20.3 edgeR_3.0.4
[13] limma_3.14.3         DESeq_1.10.1          lattice_0.20-10
[16] locfit_1.5-8         Biobase_2.18.0        Rsamtools_1.10.2
[19] Biostrings_2.26.2    rtracklayer_1.18.1    GenomicRanges_1.10.5
[22] IRanges_1.16.4       BiocGenerics_0.4.0

loaded via a namespace (and not attached):
[1] annotate_1.36.0        AnnotationForge_1.0.3 biomaRt_2.14.0
[4] bitops_1.0-4.2        BSgenome_1.26.1      genefilter_1.40.0
[7] geneplotter_1.36.0    grid_2.15.2          GSEABase_1.20.1
[10] hwriter_1.3           parallel_2.15.2      plyr_1.7.1
[13] RBGL_1.34.0           RColorBrewer_1.0-5   RCurl_1.95-3
[16] splines_2.15.2        statmod_1.4.16       stats4_2.15.2
[19] stringr_0.6.1         survival_2.36-14     tools_2.15.2
[22] XML_3.95-0.1          zlibbioc_1.4.0
```