

Mapping, Alignment and SNP Calling

Heng Li

Broad Institute

MPG Next Gen Workshop 2011

Outline

1 Mapping

- Messages from the 1000 Genomes Project
- A race in throughput

2 Alignment

- Mapping vs. alignment
- Fixing wrong alignments

3 SNP calling

- Single-sample SNP calling
- Multi-sample SNP calling

Outline

1 Mapping

- Messages from the 1000 Genomes Project
- A race in throughput

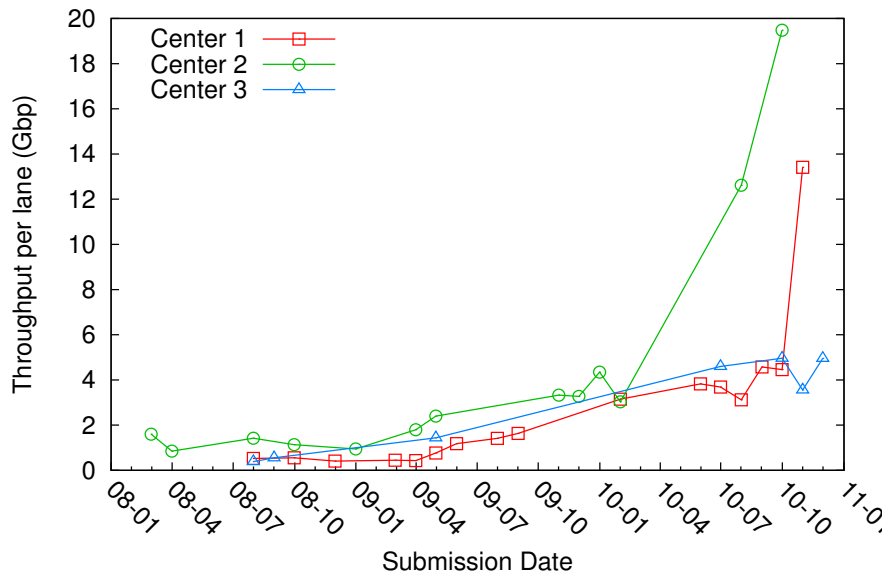
2 Alignment


- Mapping vs. alignment
- Fixing wrong alignments

3 SNP calling

- Single-sample SNP calling
- Multi-sample SNP calling


Submitted Illumina data from the 1000 Genomes Project






Illumina sequencing


- >20X increased throughput in 3 years
- ~20Gbp raw sequences per machine day at present



Published general-purpose NGS mappers



Published general-purpose NGS mappers



Convergence in mapping algorithms

- Recommended mappers for *variant calling*:
 - ▶ Illumina: BWA, Eland2, Novoalign, Stampy
 - ▶ SOLiD: Bfast, BioScope
 - ▶ 454: SSAHA2, gsMapper, BWA-SW
- Modern short-read mappers are faster than image analysis and base calling.
 - ▶ No need for further speed improvements for short reads
 - ▶ Long-read and multi-reference alignments still pose challenges

Outline

1 Mapping

- Messages from the 1000 Genomes Project
- A race in throughput


2 Alignment

- Mapping vs. alignment
- Fixing wrong alignments

3 SNP calling

- Single-sample SNP calling
- Multi-sample SNP calling

Ungapped mappers perform badly for SNP calling



A typical SNP caller sees...

```
9   t   ttt
10  a   aaaC
11  a   aaaaa
12  a   aaaaaa
13  a   aaaaaa
14  c   cccTTT
15  a   aaaaaa
16  a   aaaaaa
17  t   AAtttt
18  t   tttttt
19  a   aaaaaa
20  a   aaaaaa
21  g   Tgggg
```

The alignment looks like...

			coord	12345678901234	5678901234567890123456
9	t	ttt	ref	aggttttataaaac----	aattaagtctacagagcaacta
10	a	aaaC	sample	aggttttataaaacAAAT	aattaagtctacagagcaacta
11	a	aaaaa	read1	aggttttataaaac	<u>aa</u> Ataa
12	a	aaaaaa	read2	ggttttataaaac	<u>aa</u> AtaaTt
13	a	aaaaaa	read3	ttataaaacAAAT	aattaagtctaca
14	c	cccTTT	read4	CaaaT	aattaagtctacagagcaac
15	a	aaaaaa	read5	<u>aa</u> T	aattaagtctacagagcaact
16	a	aaaaaa	read6	<u>T</u>	aattaagtctacagagcaacta
17	t	AAtttt			
18	t	tttttt			
19	a	aaaaaa			
20	a	aaaaaa			
21	g	Tgggg			

But what is really happening is...

		coord	12345678901234	5678901234567890123456
9	t	ttt	ref	aggttttataaaac----aattaagtctacagagcaacta
10	a	aaaC	sample	aggttttataaaacAAATaattaagtctacagagcaacta
11	a	aaaaa	read1	aggttttataaaac aaAtaa
12	a	aaaaaa	read2	ggttttataaaac aaAtaaTt
13	a	aaaaaa	read3	ttataaaacAAATaattaagtctaca
14	c	cccTTT	read4	CaaaTaattaagtctacagagcaac
15	a	aaaaaa	read5	aaTaattaagtctacagagcaact
16	a	aaaaaa	read6	Taattaagtctacagagcaacta
17	t	AAtttt	read1	aggttttataaaacaaataa
18	t	tttttt	read2	ggttttataaaacaaataatt
19	a	aaaaaa	read3	ttataaaacaaataattaagtctaca
20	a	aaaaaa	read4	caataaattaagtctacagagcaac
21	g	Tgggg	read5	aataaattaagtctacagagcaact
			read6	taattaagtctacagagcaacta

Mapping vs. alignment

Mapping

- A mapping is the region where a read sequence is placed.
- A mapping is regarded to be correct if it overlaps the true region.

Alignment

- An alignment is the detailed placement of each base in a read.
- An alignment is regarded to be correct only if each base is placed correctly.

The problem

- A read mapper is fairly good at mapping, may not be good at alignment.
- This is because the true alignment minimizes differences between reads, but the read mapper only sees the reference.

Mapping vs. alignment

Mapping

- A mapping is the region where a read sequence is placed.
- A mapping is regarded to be correct if it overlaps the true region.

Alignment

- An alignment is the detailed placement of each base in a read.
- An alignment is regarded to be correct only if each base is placed correctly.

The problem

- A read mapper is fairly good at mapping, may not be good at alignment.
- This is because the true alignment minimizes differences between reads, but the read mapper only sees the reference.

Fixing wrong alignments

Multi-sequence realignment

- Perform multi-seq alignment to minimize differences between reads.
- Effective for long gaps.

Base Alignment Quality (BAQ)

- Measure the probability of a read base being misaligned
 - ▶ BAQ is low if the read base is aligned to a different reference base in a suboptimal alignment.
 - ▶ Bases with low BAQ ignored or downweighted in SNP calling.
- Effective even if no reads are mapped with gaps.
- Work by traversing all the possible alignment between the read and the reference.
- Computed efficiently with an HMM.

Fixing wrong alignments

Multi-sequence realignment

- Perform multi-seq alignment to minimize differences between reads.
- Effective for long gaps.

Base Alignment Quality (BAQ)

- Measure the probability of a read base being misaligned
 - ▶ BAQ is low if the read base is aligned to a different reference base in a suboptimal alignment.
 - ▶ Bases with low BAQ ignored or downweighted in SNP calling.
- Effective even if no reads are mapped with gaps.
- Work by traversing all the possible alignment between the read and the reference.
- Computed efficiently with an HMM.

Evaluation on simulated data

GATKrealn	BAQ	FNR	# false SNPs
No	No	7.3%	116
Yes	No	7.6%	4
No	Yes	8.3%	2
Yes	Yes	8.3%	0

- No filtering applied except a quality cutoff
- BAQ and multi-sequence realignment complement each other:
 - ▶ BAQ is less effective given long gaps.
 - ▶ The current realignment algorithm is less effective if no reads are mapped with gaps.

Outline

1 Mapping

- Messages from the 1000 Genomes Project
- A race in throughput

2 Alignment

- Mapping vs. alignment
- Fixing wrong alignments

3 SNP calling

- Single-sample SNP calling
- Multi-sample SNP calling

Single-sample Bayesian caller: a toy example

Input

Reference is C, observing 4C and 2T, all with base quality 30.

Likelihood of data

- $P(D|CC) = \Pr\{\text{two Q30 errors}\} = 10^{-(30+30)/10} = 10^{-6}$
- $P(D|TT) = \Pr\{\text{four Q30 errors}\} = 10^{-(30*4)/10} = 10^{-12}$
- $P(D|CT) = \Pr\{\text{sample 6 reads from 2 chr}\} = 1/2^6 = 1.56 \times 10^{-2}$

Posterior

- Prior: $P(CC) = 0.9985$, $P(CT) = 0.001$ and $P(TT) = 0.0005$

$$P(CC|D) = \frac{P(D|CC)P(CC)}{P(D|CC)P(CC) + P(D|CT)P(CT) + P(D|TT)P(TT)}$$

- Get: $P(CC|D) = 0.06$, $P(CT|D) = 0.94$ and $P(TT|D) = 3 \times 10^{-11}$

Single-sample Bayesian caller: a toy example

Input

Reference is C, observing 4C and 2T, all with base quality 30.

Likelihood of data

- $P(D|CC) = \Pr\{\text{two Q30 errors}\} = 10^{-(30+30)/10} = 10^{-6}$
- $P(D|TT) = \Pr\{\text{four Q30 errors}\} = 10^{-(30*4)/10} = 10^{-12}$
- $P(D|CT) = \Pr\{\text{sample 6 reads from 2 chr}\} = 1/2^6 = 1.56 \times 10^{-2}$

Posterior

- Prior: $P(CC) = 0.9985$, $P(CT) = 0.001$ and $P(TT) = 0.0005$

$$P(CC|D) = \frac{P(D|CC)P(CC)}{P(D|CC)P(CC) + P(D|CT)P(CT) + P(D|TT)P(TT)}$$

- Get: $P(CC|D) = 0.06$, $P(CT|D) = 0.94$ and $P(TT|D) = 3 \times 10^{-11}$

Single-sample Bayesian caller: a toy example

Input

Reference is C, observing 4C and 2T, all with base quality 30.

Likelihood of data

- $P(D|CC) = \Pr\{\text{two Q30 errors}\} = 10^{-(30+30)/10} = 10^{-6}$
- $P(D|TT) = \Pr\{\text{four Q30 errors}\} = 10^{-(30*4)/10} = 10^{-12}$
- $P(D|CT) = \Pr\{\text{sample 6 reads from 2 chr}\} = 1/2^6 = 1.56 \times 10^{-2}$

Posterior

- Prior: $P(CC) = 0.9985$, $P(CT) = 0.001$ and $P(TT) = 0.0005$

$$P(CC|D) = \frac{P(D|CC)P(CC)}{P(D|CC)P(CC) + P(D|CT)P(CT) + P(D|TT)P(TT)}$$

- Get: $P(CC|D) = 0.06$, $P(CT|D) = 0.94$ and $P(TT|D) = 3 \times 10^{-11}$

Multi-sample Bayesian caller: an overview

- Similar to single-sample calling except replacing the individual genotype with the genotype configuration of multiple samples.
- Math magic to accelerate computation.
- Allele frequency estimated at the same time.

Multi-sample vs. pooled SNP calling

An example

- 1 sample covered by 3 Q20 C bases (1% error rate); 99 samples covered by 297 Q20 T bases.
- Very unlikely for 3 errors appear in one sample.
- Without sample information, the 3 C look like perfect sequencing errors.

Combining pooling and barcoding

- Pool less than 100 samples together, barcode each pool and then sequence.

Acknowledgements

- 1000 Genomes Project analyses group
- Mark Depristo and the GSA group at Broad
- SAMtools/Picard users
- Altshuler/Daly lab and Reich lab

Thank You